

I/O OPERATIONS

KONRAD HINSEN

CENTRE DE BIOPHYSIQUE MOLÉCULAIRE (ORLÉANS)

AND

SYNCHROTRON SOLEIL (ST AUBIN)

TEXT AND BINARY FILES

There are two big families of file format:

- text files store data as human-readable text
- binary files store data as compactly as possible

Word-processor files
are often binary files!

Advantages of text files:

- you can just look at them to see what's inside
- there are thousands of tools that work with them

Advantages of binary files:

- they typically use less space
- they are handled faster by the computer

We will look at binary
files in part 2!

Convenience/performance tradeoff

Example: in 2 bytes, you can store integers
up to $2^{16} = 65\,536$ in a binary file,
but only up to 99 in a text file.

ASCII

Codes 20 to 7E (in decimal: 32 to 126) are standard characters

Codes 00 to 1F (in decimal: 0 to 31) are control characters,
originally used in the communications protocol for teletype printers

Code 7F (in binary: 1111111) was used to erase data on punched
tapes by overwriting.

ASCII Code Chart

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | A | B | C | D | E | F |
|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|----|-----|-----|----|----|----|-----|
| 0 | NUL | SOH | STX | ETX | EOT | ENQ | ACK | BEL | BS | HT | LF | VT | FF | CR | SO | SI |
| 1 | DLE | DC1 | DC2 | DC3 | DC4 | NAK | SYN | ETB | CAN | EM | SUB | ESC | FS | GS | RS | US |
| 2 | | ! | " | # | \$ | % | & | ' | (|) | * | + | , | - | . | / |
| 3 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | : | ; | < | = | > | ? |
| 4 | @ | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O |
| 5 | P | Q | R | S | T | U | V | W | X | Y | Z | [| \ |] | ^ | _ |
| 6 | ` | a | b | c | d | e | f | g | h | i | j | k | l | m | n | o |
| 7 | p | q | r | s | t | u | v | w | x | y | z | { | | } | ~ | DEL |

ASCII CONTROL CODES

Most ASCII control codes were designed for teletypes and are obsolete.

Two control codes are frequent in text files:

- LF (line feed, code 10)
- CR (carriage return, code 13)

One or the other, or both (in the order CR, LF) are used to indicate the end of a line. Unix uses LF, MacOS (before MacOS X) used CR, DOS and Windows use CR+LF.

More rarely used control codes:

- BS (backspace, code 8)
- HT (horizontal tab, code 9)
- FF (feed forward, code 12)

UNICODE

The Unicode standard aims at covering *all* the world's languages in a single file. This is not possible with just 256 values for a byte. Unicode currently defines 1 114 112 characters, each of which requires 21 bits of storage.

The first 127 codes are equal to ASCII.

Unicode is most commonly stored in files using UTF-8 encoding (Unicode Transformation Format using 8 bits). In UTF-8, the first 127 characters (ASCII) are stored as a single byte. Other characters are stored as a sequence of two or three bytes.

Every valid ASCII file is a valid UTF-8 file.

UTF-8 is on its way to becoming the most widely used character encoding system, in particular on the World-Wide Web.

BINARY REPRESENTATIONS FOR NUMBERS

Integers:

- each number is represented by N bits (N=8, 16, 32, 64)
- *unsigned* integers: range from 0 to 2^N-1
- *signed* integers: range from -2^{N-1} to $2^{N-1}-1$

Order of the individual bytes in memory or in a file:

- highest-value byte first: “big-endian”
- lowest-value byte first: “little-endian”

| bits | unsigned | signed |
|------|----------|--------|
| 00 | 0 | 0 |
| 01 | 1 | 1 |
| 10 | 2 | -2 |
| 11 | 3 | -1 |

BINARY REPRESENTATIONS FOR NUMBERS

Floating-point numbers (IEEE 754):

- general form of a float: $(-1)^s \cdot c \cdot b^q$
- *sign* s : 0 or 1 (one bit)
- *base* b : 2 or 10
- *significand* c : integer in the range 0 to b^p-1
- *precision* p : positive integer
- *exponent* q : integer with range $1-e_{\max} \leq q+p-1 \leq e_{\max}$

Most common base-2 varieties:

- single-precision: $p = 24$, $e_{\max} = 127$, stored in 4 bytes
- double-precision: $p = 53$, $e_{\max} = 1023$, stored in 8 bytes

Special values:

- NaN (not a number): result of undefined operations, e.g. logarithm of a negative number
- positive and negative infinity

BINARY FILES

Traditionally:

- each program defines its own binary file formats, arranged for efficiency and the programmer's convenience
- binary formats tend to be undocumented
- some binary formats are platform-dependent (size of integers, byte order)
- data often becomes unusable over time, as platforms and programs evolve

Evolution towards standardized binary formats in science:

- platform-independent binary representations for integers (XDR) and floats (IEEE) since the 1980s
- standardized special-purpose formats (e.g. FITS)
- standardized general-purpose formats: CDF, netCDF, HDF

NON-STANDARD BINARY FILES

Avoid them if you can!

If you can't:

- get the documentation if there is one
- get the source code of the program otherwise
- get a collection of example files
- use `od` to look at the examples, and check they agree with the documentation
- make sure you know what each bit means!
- find out how your favorite programming languages handles binary file I/O

Examples

- Python: use modules `struct` and `ctypes`
- C/C++: `fopen()` / `fread()` / `fwrite()` / `fclose()`
- Fortran: unformatted I/O

FILE FORMATS

WHAT'S A FILE FORMAT?

In theory: a precise definition of

- what data is stored
- how it is represented
- in which order items are stored

In practice: a more or less precise definition

Ideally, a file format should allow *validation*: running a program that decides if a given file implements the format correctly or not. Validation requires a precise definition of the file format.

Good file formats are important for:

- exchanging data between different programs
- archiving data for long-term use

WHAT'S A DATA MODEL?

A definition of how complex data is represented in terms of basic data (numbers, text, tables, ...). A data model can be realized in many different concrete file formats.

Example: the PDB data model

- a list of atoms with for each atom
 - ♦ a residue name and residue number
 - ♦ an atom name
 - ♦ the chemical element symbol
 - ♦ a position (x, y, z coordinates)
 - ♦ the uncertainty of the position (B-factor or ADP)
- nature of the molecules
- experimental conditions
- ...

File formats: PDB, mmCIF, PDBML

PROPRIETARY FILE FORMATS

Proprietary file formats are undocumented file formats used exclusively by a set of programs and inaccessible to others.

Some vendors use proprietary formats out of laziness, but usually the goal is to force you to use their products exclusively.

Don't use proprietary file formats for scientific data.

It's you who is responsible for your data, so you must be able to control it to the last bit if necessary.

MARKUP LANGUAGES

Convention for indicating *semantic* or *presentation* features in human-readable text.

Examples: HTML, TeX, Markdown, reST, ...

Application: text that is processed by computer programs

Examples: Web pages, scientific articles, software documentation, comments in scientific data sets, ...

FULLY STRUCTURED DATA

Fully structured data has a completely predefined structure. The structure is defined by the file format, the file contains only the data.

Examples: tables, arrays, database records, ...

Example file formats: CSV, PDB, FASTA

Fully structured data is the easiest to handle in computer programs, but also the least flexible.

SEMI-STRUCTURED DATA

Semi-structured data *contains* tags that define its own structure. There is a general convention that specifies how to represent basic data items (numbers, text, ...), and for each format a specific convention that specifies the valid tags and the rules for combining tagged items.

General conventions: XML, SGML, s-expressions,
JSON, ...

Specific file formats: PDBML (XML), HTML (SGML),
Lisp (s-expressions), ...

Semi-structured data is much more flexible than fully structured data and allows in particular future extensions. Most recently defined file formats are based on semi-structured formats, in particular XML.

XML: EXTENSIBLE MARKUP LANGUAGE

Published in 1996 as a simplified version of SGML (Standard Generalized Markup Language). Became rapidly popular for Web technologies.

An XML file consists of Unicode characters.

Tags are surrounded by < > and can contain attributes.

Elements are delimited by a start tag and a matching end tag:

```
<pdb-id> 1IEE </pdb-id>
```

Empty elements use special tags:

```

```

A specific XML-based format is defined by a *Document Type Definition* (DTD) or a *Schema* in order to allow validation.

Scientific XML formats: PDBML (PDB), CML (chemistry),
SBML (systems biology), MathML (maths),
XSIL (tables and arrays), ...

FITS

FITS IN PYTHON

- Formerly PyFITS, now part of astropy
- Tutorial at
<http://www.astropython.org/tutorial/2010/10/PyFITS-FITS-files-in-Python>

HDF5

(HIERARCHICAL DATA FORMAT 5)

HDF5 OVERVIEW

Data model: **arrays and tables in a file**

- Widely used in different communities
- Basis of other formats (NeXus, netCDF4)
- Works on all current computing platforms
- Portable data files of unlimited size
- Many tools available
- Library written in C
- Bindings for C++, Fortran 90, and Java
- Supports parallel I/O

<http://www.hdfgroup.org/HDF5/>

Third-party bindings:

Python (h5py, PyTables), Matlab, IDL, Mathematica

STRUCTURE OF A HDF5 FILE

Dataset: multidimensional array of data elements plus metadata (name, description, time stamps, ...)

Group: a named grouping of datasets and other groups

Node: a group or a dataset

HDF5 file: contains the “root group” of a group hierarchy

Analogy to a file system: datasets \leftrightarrow files, groups \leftrightarrow directories.

Path notation to specify groups and datasets:

/group1/group2/dataset

DATASETS

Name: any text string

Datatype:

- *atomic*: integer, float, string of various sizes
- *compound*:
 - multidimensional array, elements of any datatype
 - variable length 1D array, elements of any datatype
 - record with named fields of any datatype

Compounds can be recursive: a record with a field that is an array of records whose fields are integers...

Dataspace: describes the arrangement of elements

- *scalar*: one element
- *simple*: multidimensional array of elements

Storage layout: contiguous, chunked, compact

ATTRIBUTES

Small datasets attached to **primary datasets** or to **groups**.
Typically used to store **metadata**.

Name: any text string

Datatype: as for datasets

Dataspace: as for datasets

Storage layout: always compact

Restrictions:

- can't have attributes
- no partial I/O

HDF5 TOOLS

Command line tools:

- **h5ls:** show names of groups and datasets in an HDF5 file
- **h5dump:** show contents of datasets
- **h5diff:** compare datasets and files
- **h5import:** import data from text and platform-specific binary files into an HDF5 file

GUI browser and editor: hdfview

HDF5 LIBRARY

Sublibraries by function category:

- H5: general library functions
- H5F: files
- H5D: datasets
- H5T: datatypes
- H5S: dataspace
- H5A: attributes

plus many more!

OO-like structure:

- **open** (file, dataset, datatype, ...) creates an in-memory data structure and returns a unique handle to it
- **close** deallocates the data structure and makes the handle invalid

HDF5 DOCUMENTATION

Introduction:

<http://www.hdfgroup.org/HDF5/doc/H5.intro.html>

User's Guide: (more detailed)

<http://www.hdfgroup.org/HDF5/doc/UG/>

Function reference:

http://www.hdfgroup.org/HDF5/doc/RM/RM_H5Front.html

Examples:

<http://www.hdfgroup.org/ftp/HDF5/examples/examples-by-api/>

HDF5 IN PYTHON

PyTables:

<http://www.pytables.org/>

database of tables stored in HDF5 files

h5py:

<http://h5py.org/>

closer to the HDF5 C/Fortran interface

arrays in a file

HDF5 IN PYTHON: WRITING

```
import h5py
import numpy as np

h5file = h5py.File('test.h5', 'w')
h5file.attrs['version'] = 42

foo = h5file.create_dataset('foo', (20, 3), dtype=np.float64)
foo[:, :] = 0.
foo[0, :] = [42., 42., 42.]
foo[:, 1] = 1.
foo.attrs['units'] = "arbitrary"

bar = h5file.create_dataset('bar', (0, 20), dtype=np.int32,
                           chunks=(10, 10), maxshape=(None, 20))
for i in range(10):
    bar.resize((i+1, 20))
    bar[i, :] = i*np.ones((1, 20), np.int)

h5file.close()
```


HDF5 IN PYTHON: READING

```
import h5py
import numpy as np

h5file = h5py.File('test.h5', 'r')

for item in h5file:
    print item

foo = h5file['foo'][...]
foo_units = h5file['foo'].attrs['units']
print foo[0]

h5file.close()
```


HDF5 IN C: WRITING DATA

```
#include "hdf5.h"
```

```
#define FILE          "demo.h5"
```

```
#define DATASETNAME "IntArray"
```

```
#define NX          5
```

```
#define NY          6
```

```
#define RANK        2
```

```
int main (void) {  
    hid_t      file, dataset;  
    hid_t      datatype, dataspace;  
    hsize_t    dims[2];  
    herr_t     status;  
    int        data[NX][NY];  
    int        i, j;  
  
    for (j = 0; j < NX; j++) {  
        for (i = 0; i < NY; i++)  
            data[j][i] = i + j;  
    }
```

```
    file = H5Fcreate(FILE, H5F_ACC_TRUNC, H5P_DEFAULT, H5P_DEFAULT);
```

```
    dims[0] = NX;
```

```
    dims[1] = NY;
```

```
    dataspace = H5Screate_simple(RANK, dims, NULL);
```

```
    datatype = H5Tcopy(H5T_NATIVE_INT);
```

```
    status = H5Tset_order(datatype, H5T_ORDER_LE);
```

```
    dataset = H5Dcreate(file, DATASETNAME, datatype, dataspace,  
                        H5P_DEFAULT);
```

```
    status = H5Dwrite(dataset, H5T_NATIVE_INT, H5S_ALL, H5S_ALL,  
                      H5P_DEFAULT, data);
```

```
    H5Sclose(dataspace);
```

```
    H5Tclose(datatype);
```

```
    H5Dclose(dataset);
```

```
    H5Fclose(file);
```

```
    return 0;
```

```
}
```


HDF5 IN FORTRAN 90: WRITING DATA

```
PROGRAM main
```

```
USE HDF5
```

```
IMPLICIT NONE
```

```
CHARACTER(LEN=14), PARAMETER :: filename = "demo.h5"
```

```
CHARACTER(LEN=3) , PARAMETER :: dataset = "DS1"
```

```
INTEGER          , PARAMETER :: dim0      = 4
```

```
INTEGER          , PARAMETER :: dim1      = 7
```

```
INTEGER :: hdferr
```

```
INTEGER(HID_T) :: file, space, dset
```

```
INTEGER(HSIZE_T), DIMENSION(1:2)          :: dims = (/dim0, dim1/)
```

```
INTEGER          , DIMENSION(1:dim0,1:dim1) :: wdata
```

```
INTEGER :: i, j
```

```
CALL h5open_f(hdferr)
```

```
DO i = 1, dim0
```

```
    DO j = 1, dim1
```

```
        wdata(i,j) = (i-1)*(j-1)-(j-1)
```

```
    ENDDO
```

```
ENDDO
```

```
CALL h5fcreate_f(filename, H5F_ACC_TRUNC_F,  
                  file, hdferr)
```

```
CALL h5screate_simple_f(2, dims, space, hdferr)
```

```
CALL h5dcreate_f(file, dataset, H5T_STD_I32LE,  
                 space, dset, hdferr)
```

```
CALL h5dwrite_f(dset, H5T_NATIVE_INTEGER,  
               wdata, dims, hdferr)
```

```
CALL h5dclose_f(dset , hdferr)
```

```
CALL h5sclose_f(space, hdferr)
```

```
CALL h5fclose_f(file , hdferr)
```

```
END PROGRAM main
```


EXERCISES

EXPLORING A NEXUS FILE

The NeXus file format is a set of conventions for storing experimental data (X-ray diffraction and neutron scattering experiments) in a HDF5 file.

- 1) Inspect the example file `Co_edge_CnrX_2008-09-17_22-12-12.nxs` using HDFView and/or the HDF5 command line tools.
(Thanks to the Proxima 1 team at Synchrotron SOLEIL for providing this file!)
- 2) Which user recorded the data? Find his name and telephone number.
- 3) Plot the `fitted_escan_curves: raw_data` vs. `x-ray_energy`, using whatever programs you are familiar with.

EXPLORING METEOROLOGICAL DATA

Inspect the example file

```
OMI.L2.CloudOMCLDO2Strip200kmAlongCloudSat.2011.06.22.050738Z.v003.he5
```

(that's the original name as obtained from the NASA Earth Observatory!)

What can you say about the data in that file with no additional information?

HDF5 IN PYTHON

Write a Python script to explore the contents of the NeXus example file.
Make it print out as much useful information as you can find.